


Supervised Learning

Neural Networks 2

Prof. Dr. Johannes Jurgovsky

Recap: Logistic Regression

- A binary classifier
- Output: linear function of input +
(non-linear) activation function (logistic function)
 $\hat{y} = g(\sum w_i x_i)$
 \hat{y} is interpreted as $p(y = 1 | x)$






Multi-Layer Perceptron (aka Neural Network)

- Try it here: <http://playground.tensorflow.org>

Multi-Layer Perceptron (MLP)

- Neurons are arranged in layers
- Layer n is (fully) connected with layer $n+1$
 - no connections within layer
 - no connections to any other layers
 - no feedback
- Information flow from one layer to the next: feed-forward
- network has no internal state
- number of input/output units is problem dependent
- number of hidden units determined by developer



MLP – Properties


A MLP can

- compute any Boolean function (AND, OR, XOR, ...)
- approximate any non-linear function
 - a 2-layer MLP with a finite number of hidden neurons suffices; however, a lot of neurons may be required in that hidden layer
- define arbitrary class boundaries
- be trained in a supervised fashion using a sample set (Error Backpropagation)
- be adapted to binary classification, multi-class classification and regression




Activation Functions (classic)

Step function/Threshold



not differentiable


Sigmoid/logistic function



$$g(x) = \frac{1}{1 + e^{-x}}$$

$$g'(x) = g(x)(1 - g(x))$$

Sigmoid/tanh



$$g(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$


$$g'(x) = 1 - g^2(x)$$

Changing the bias $b = w_0$ moves threshold



Activation function – (Leaky) ReLU / Softplus


ReLU



$$g(x) = \max(0, x)$$

$$g'(x) = \begin{cases} 0 & x \leq 0 \\ 1 & x > 0 \end{cases}$$


Leaky ReLU



$$g(x) = \begin{cases} 0.01x & x \leq 0 \\ x & x > 0 \end{cases}$$

$$g'(x) = \begin{cases} 0.01 & x \leq 0 \\ 1 & x > 0 \end{cases}$$

Softplus



$$g(x) = \ln(1 + e^x)$$

$$g'(x) = \frac{1}{1 + e^{-x}}$$

- Ramp function
- ReLU = Rectified Linear Unit
- by now the most common activation functions in deep neural networks (for inner neurons)
- Variant Softplus: smooth approximation of ReLU


Activation Function – Softmax

- A normalized exponential function $g(x)_j = \frac{e^{x_j}}{\sum_i^C e^{x_i}}$
- Smooth approximation of MAX-Function
 - lifts high values and suppresses low values
- Used in output neurons
- Acts like a probability mass function over C classes
- Example
 - in 1 2 3 4 1 2 3
 - out 0,024 0,064 0,175 0,475 0,024 0,064 0,175

Training

Training = Determine weights

1. Feed training examples through the network (forward pass)
2. Error-Backpropagation (backward pass)



Deep learning

- = representation learning = learn feature extraction
- use many layers of neurons
- Example: AlexNet (image recognition)
 - Convolutional Neural Network (CNN) → Computer Vision (winter term)
 - 8 layers (not all of these fully connected)
 - 650,000 neurons
 - 60 million parameters
 - trained on 1.2 million images
- Example: ResNet – 50 to 150 layers deep

MLP – Training (Error-Backpropagation)

- Feed training examples through network
- Compute error for each example n and each output neuron j :

$$e_j(n) = y_j(n) - a_j(n)$$
(y : desired output, $a = \hat{y}$ actual output)

- Minimize total error of all output neurons: $\varepsilon(n) = \frac{1}{2} \sum_j e_j^2(n)$
 - here: Mean Squared Error (MSE)
 - using non-linear optimization (e.g. gradient descent; requires partial derivatives of error function)
 - result: *local* optimum
 - the required update of a weight between neuron j of layer l and neuron i of layer $l - 1$ is:

$$\Delta w_{ji}(n) = -\alpha \frac{\partial \varepsilon(n)}{\partial z_j(n)} a_i(n)$$

where $z_j(n)$ is the weighted sum of neuron inputs and α is the learning rate

- Partial derivative for neurons in output layer:

$$\frac{\partial \varepsilon(n)}{\partial z_j(n)} = -e_j(n) g'(z_j(n))$$

- Partial derivative for hidden neurons in layer l :

$$\frac{\partial \varepsilon(n)}{\partial z_j(n)} = g'(z_j(n)) \sum_k \frac{\partial \varepsilon(n)}{\partial z_k(n)} w_{kj}(n)$$

where k iterates over all neurons of layer $l + 1$

Normalization / Weight Initialization

- Input data:
Normalize, such that all values are of similar magnitude
- Weight Initialization:
 - Initialization of all weights with zeros or a single constant is bad :
Neurons in hidden layers collapse to a single neuron
(because of symmetry)
 - Random initialization: common heuristics, e.g.
 - uniformly from the interval $[-0,01; +0,01]$
 - uniformly from the interval $[-\frac{1}{\sqrt{n}}; +\frac{1}{\sqrt{n}}]$, where n = number of neurons of previous layer
 - for deep neural networks: uniformly from the interval $[-\frac{\sqrt{6}}{\sqrt{m+n}}; +\frac{\sqrt{6}}{\sqrt{m+n}}]$, where m = #neurons of previous layer, n = #neurons of subsequent layer
- Initialization of bias: Zero

Notes

- Advantages:
 - parallel computation within layer possible
 - can be formulated as matrix multiplications
→ well suited for GPUs
- Disadvantages:
 - Number of hidden layers and number of neurons per hidden layer is a design decision
 - theoretical minimum of 2/3 layers
(but more may be better, e.g., faster convergence, less neurons)
 - there are presently no well-founded results on how to choose these hyper-parameters → needs to be done experimentally
 - high computation times for training as well as classification
 - non-linear optimization guarantees only local minimum
- Try it here: <http://playground.tensorflow.org>

Learning Rate

- Learning Rate too high
 - Optimization gets stuck on plateau
 - or even diverges
- Learning Rate too small
 - slow convergence
- Start with high values, e.g. $\alpha = 0.1$ or even higher
- If diverging: Decrease learning rate
- Typical learning rates are (depending on optimization method) 0.01 or 0.001

Optimization Method – Mini Batch SGD

- SGD = Stochastic Gradient Descent (SGD)
- Most used
- Faster than standard gradient descent
- For each iteration k :
 - select a small set of training examples randomly (mini-batch)
 - compute gradients
 - update weights
 - decrease the learning rate linearly:
 $\alpha_k = (1 - \beta)\alpha_0 + \beta\alpha_\tau$ where $\beta = \frac{k}{\tau}$
learning rate stays constant after iteration τ

Optimization Method – Momentum

- accelerated trainings for
 - high curvature
 - small, but consistent gradients
 - noisy gradients
- Idea: Store gradients from previous updates
 - as floating mean
 - influence of older values decreases exponentially

- Accumulate gradients

$$m_t = \gamma m_{t-1} + \alpha \delta d$$

- Update parameters

$$w_t = w_{t-1} - m_t$$

- m : moment, α : learning rate, δ : partial derivative in corresponding layer, d : output of neurons of previous layer, w : weights
 γ : dampening factor (e.g. 0.9)

- Nesterov-Moments

- gradient is computed after weights have been updated in current iteration
- basically, a correction term to account for changes in gradient

Optimization Method – Adaptive Learning Rate

- selecting the "correct" learning rate is
 - critical – high impact on result
 - difficult
- slightly less so when using moments
 - but at the cost of an additional parameter
- adaptive learning rate:
 - determine learning rate separately for each weight update (batch) based on a global learning rate parameter

Optimization Method – Adaptive Learning Rate

- AdaGrad (2011)
 - changes learning rate proportional to square root of past squared gradients
 - works well for some deep networks, but not for others
- RMSProp (2012)
 - modified AdaGrad, working better with neural networks
 - uses gradient accumulation with exponential dampening
 - there is a variant in combination with Nesterov-Moments
 - widely used for neural networks
- Adam (2014)
 - combination of RMSProp with moments
 - uses also second order moments
- in general, the training result is highly dependent on parameter settings

Regularization – Dropout

- Regularization: avoid extreme values for weights
- one possibility: Dropout
 - in each iteration: set output of some randomly selected neurons to zero (as if they had been removed)
 - typical value: 0.5 (= 50% probability for removal)
- Advantage:
 - forces network to learn redundant representations of classes
 - prohibits that different neurons concentrate on the same features
- behaves like training a large set of network models that share parameters (Bagging)

Objective Function / Loss Function

- Mean Squared Error (MSE):
see Error-Backpropagation
- Cross Entropy (or Log Loss)
 - for comparing two probability densities
 - Computation (for a single sample): $-\sum_{i=0}^{m-1} y_i \ln p_i$
 m : #classes,
 y_i : desired output (typically 0 or 1),
 p_i : actual output
- Binary Cross Entropy
 - Cross Entropy for two classes
 - Computation simplifies to: $-(y \ln p + (1 - y) \ln (1 - p))$

Activation/Loss Function – Classification

- Two classes
 - use a single neuron in output layer
 - activation output layer: sigmoid (logistic)
 - activation hidden layers: ReLU
 - loss function: binary cross entropy
- Multiple disjoint classes
 - use one neuron per class (1-out-of-n coding, one-hot) in output layer
 - activation output layer: Softmax
 - activation hidden layers: ReLU
 - loss function: cross entropy
- Multiple non-disjoint classes
 - use one neuron per class in output layer
 - activation output layer: sigmoid (logistic)
 - activation hidden layers: ReLU
 - loss function: binary cross entropy (summed over all output neurons)

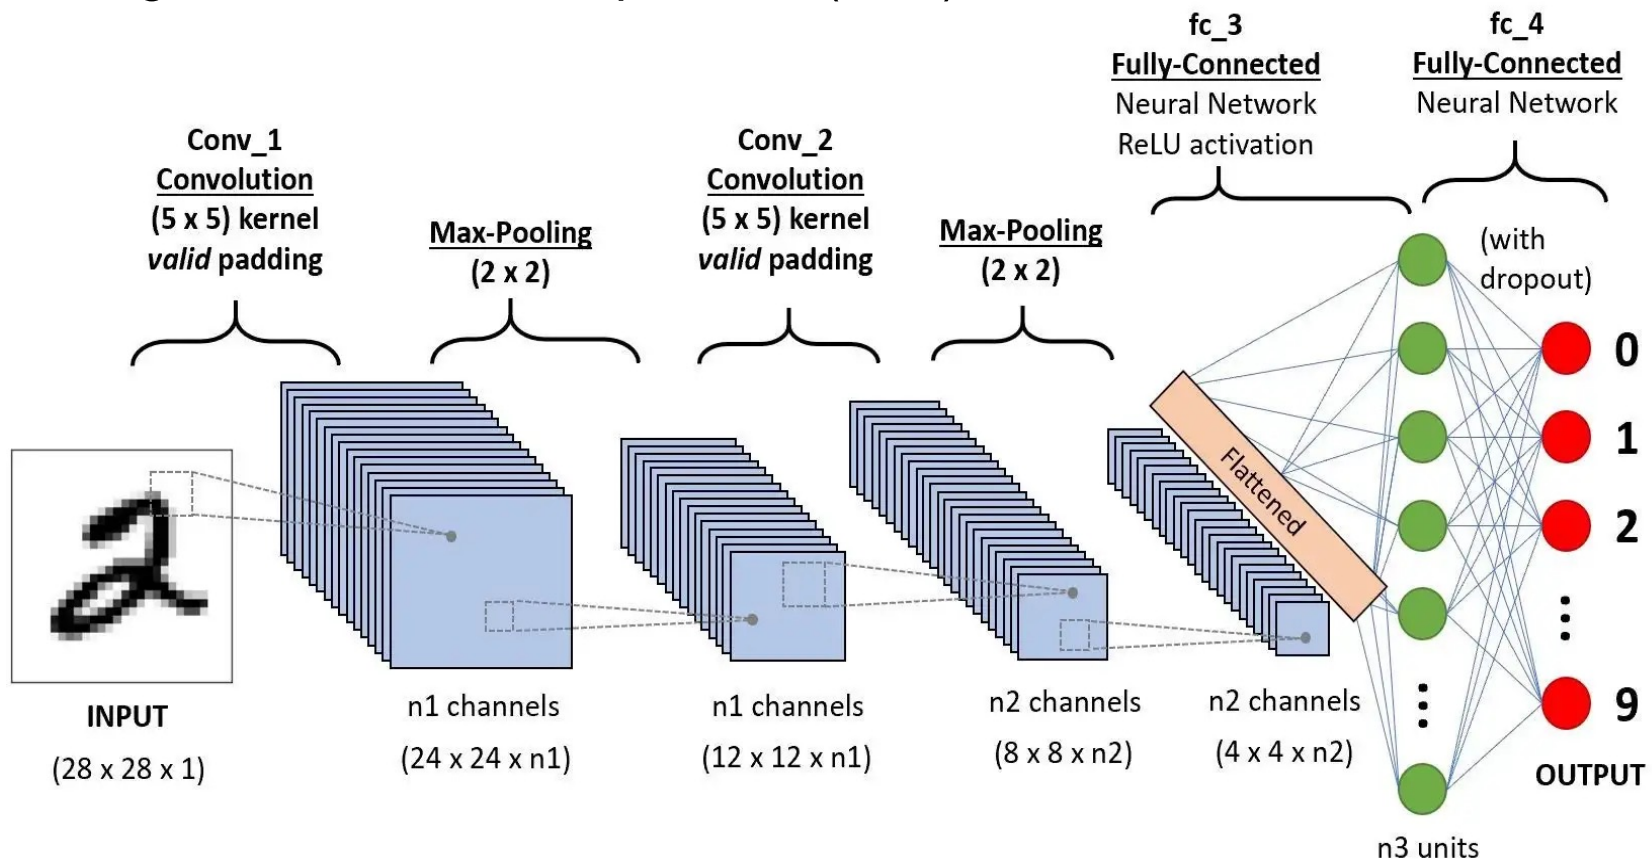
Activation/Loss Function – Regression

- instead of discrete classes: compute function value
- use a single neuron in output layer
- activation output layer: linear
(sum of weights passes through unchanged)
- activation hidden layers: ReLU
- loss function: mean squared error



Convolutional Neural Network (CNN)

- deep feed-forward networks
- not fully connected
- integrates convolution operation (filter) in network




[Source](#)


Recurrent Neural Network (RNN)

- for sequential data
- maintains internal state
- widely used: Long Short-Term Memory (LSTM) networks


one to one




one to many




many to one



many to many



many to many



[Source](#)

Neural Network APIs

- scikit-learn contains a neural network API
 - should not be used
(not intended for large scale applications, no GPU support)
- PyTorch
 - Facebook neural network API (open source)
 - <https://pytorch.org/>
- Tensorflow
 - Google Python API for neural networks (open source)
 - <https://www.tensorflow.org/>